

# Logistic Regression – Complete Beginner-to-Advanced Revision Guide

---

## 1. Introduction

### What is Logistic Regression?

Logistic Regression is a **supervised learning** algorithm used for **classification**, primarily **binary classification** (two possible outcomes).

Despite its name, logistic regression is **not a regression model for continuous outputs**. It models the **probability** that an input belongs to a particular class.

**Typical output:**  $P(y = 1 \mid x) \in [0, 1]$

---

### When to Use Logistic Regression

- Binary outcomes (Yes/No, Spam/Not spam)
  - Multiclass problems (with extensions)
  - Interpretable models
  - Probabilistic predictions required
- 

### Real-World Applications

- Medical diagnosis (disease vs no disease)
  - Credit scoring (default vs non-default)
  - Marketing (customer churn)
  - Fraud detection
  - Admission prediction
- 

## 2. Mathematical Foundations

---

### 2.1 Linear Model

We start with a **linear combination** of features:

$$z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n = \mathbf{w}^T \mathbf{x}$$

---

Where:

- $w$  = weights (parameters)
  - $x$  = input features
  - $z$  = logit (raw score)
- 

## 2.2 Why We Need the Sigmoid Function

Linear regression outputs values from  $(-\infty, +\infty)$ , but probabilities must lie in  $[0, 1]$ .

### Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

#### Properties:

- Smooth, differentiable
  - Maps real numbers  $\rightarrow$  probabilities
  - S-shaped curve
- 

## 2.3 Logistic Regression Model

Combining linear model and sigmoid:

$$P(y=1 | x) = \sigma(\mathbf{w}^T \mathbf{x})$$

## 2.4 Odds and Log-Odds

### Odds

$$\text{Odds} = \frac{P(y=1)}{1 - P(y=1)}$$

### Log-Odds (Logit)

$$\log \left( \frac{P(y=1)}{1 - P(y=1)} \right) = \mathbf{w}^T \mathbf{x}$$

**Key insight:** Logistic regression is **linear in log-odds**, not probabilities.

---

## 3. Loss Function and Training

## 3.1 Likelihood Function

For a dataset of  $m$  samples:

$$L(\mathbf{w}) = \prod_{i=1}^m P(y_i | x_i)$$

For binary labels:

$$P(y_i | x_i) = \hat{y}_i^{y_i} (1 - \hat{y}_i)^{1 - y_i}$$

## 3.2 Log-Likelihood

Taking log (numerical stability):

$$\ell(\mathbf{w}) = \sum_{i=1}^m \left[ y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i) \right]$$

## 3.3 Binary Cross-Entropy Loss

Loss to minimize:

$$J(\mathbf{w}) = -\frac{1}{m} \sum_{i=1}^m (\hat{y}_i - y_i) x$$

## 3.4 Gradient Descent

### Gradient

$$\frac{\partial J}{\partial w_j} = \frac{1}{m} \sum_{i=1}^m (\hat{y}_i - y_i) x$$

### Update Rule

$$w_j := w_j - \alpha \frac{\partial J}{\partial w_j}$$

Where:

- $\alpha$  = learning rate

## 4. Regularization

## 4.1 Why Regularization?

Prevents:

- Overfitting
  - Large coefficients
  - Poor generalization
- 

## 4.2 L2 Regularization (Ridge)

$$J(\mathbf{w}) = \text{Loss} + \frac{\lambda}{2m} \sum w_j^2$$

---

## 4.3 L1 Regularization (Lasso)

$$J(\mathbf{w}) = \text{Loss} + \frac{\lambda}{m} \sum |w_j|$$

---

**Key difference:**

- L1 → feature selection (sparse)
  - L2 → weight shrinkage
- 

## 5. Model Evaluation

---

### 5.1 Confusion Matrix

Actual \ Predicted	0	1
0	TN	FP
1	FN	TP

---

### 5.2 Metrics

#### Accuracy

$$\frac{TP + TN}{TP + TN + FP + FN}$$

## Precision

$\frac{TP}{TP + FP}$

## Recall (Sensitivity)

$\frac{TP}{TP + FN}$

## F1-Score

$2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$

---

## 5.3 ROC Curve and AUC

- ROC: TPR vs FPR
  - AUC measures ranking ability
  - Threshold-independent
- 

## 6. Interpretation of Coefficients

---

### Coefficient Meaning

$\beta_j > 0 \Rightarrow x_j \text{ increases probability}$

### Odds Ratio

$e^{\beta_j}$

-

1 increases odds - <1 decreases odds

---

## 7. Assumptions

- Binary outcome
  - Independent observations
  - Linearity in log-odds
  - No multicollinearity
  - Large sample size preferred
-

## 8. Common Pitfalls & Misconceptions

- ✗ Interpreting coefficients as linear change in probability
  - ✗ Using accuracy on imbalanced data
  - ✗ Forgetting feature scaling
  - ✗ Assuming normality of variables
  - ✗ Perfect separation causing infinite coefficients
- 

## 9. Python Implementation

---

### 9.1 From Scratch (NumPy)

```
import numpy as np

def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def train_logistic(X, y, lr=0.1, epochs=1000):
    m, n = X.shape
    w = np.zeros(n)

    for _ in range(epochs):
        z = X @ w
        y_hat = sigmoid(z)
        gradient = (1/m) * X.T @ (y_hat - y)
        w -= lr * gradient

    return w
```

---

### 9.2 Using scikit-learn

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

model = LogisticRegression(penalty='l2', solver='lbfgs')
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))
```

---

## 9.3 Decision Boundary Visualization

```
import matplotlib.pyplot as plt

plt.scatter(X[:,0], X[:,1], c=y)
plt.show()
```

---

## 10. Extensions

- Multinomial Logistic Regression (Softmax)
  - One-vs-Rest
  - Regularized Logistic Regression
  - Bayesian Logistic Regression
- 

## 11. When NOT to Use Logistic Regression

- Highly non-linear boundaries
  - Complex feature interactions
  - Very small datasets
- 

## 12. Best Practices

- Scale features
  - Tune regularization
  - Use ROC-AUC for imbalance
  - Check multicollinearity (VIF)
  - Interpret coefficients carefully
- 

## 13. Summary

Logistic Regression is:

- Simple
- Powerful

- Interpretable
- Foundational for ML understanding

Mastering it builds intuition for:

- Neural networks
  - Generalized linear models
  - Probabilistic ML
- 

End of Revision Guide