

# K-Means Clustering – Complete Beginner-to-Advanced Revision Guide

---

## 1. Introduction

### What is K-Means Clustering?

K-Means is an **unsupervised learning algorithm** used to partition a dataset into **K distinct, non-overlapping clusters** based on similarity.

Each cluster is represented by its **centroid**, and the algorithm aims to minimise **within-cluster variance**.

---

### Core Intuition

"Group points so that those within a cluster are as close as possible."

- Clusters are defined by proximity
  - Each point belongs to exactly one cluster
  - Centroids represent cluster centres
- 

### Real-World Applications

- Customer segmentation
  - Image compression
  - Market basket analysis
  - Document clustering
  - Anomaly detection (with care)
- 

## 2. Mathematical Foundations

---

### 2.1 Objective Function (Inertia)

K-Means minimises the **within-cluster sum of squares (WCSS)**:

$$\$ J = \sum_{k=1}^K \sum_{x_i \in C_k} |x_i - \mu_k|^2 \$$$

Where: -  $C_k$  = cluster k -  $\mu_k$  = centroid of cluster k

---

## 2.2 Distance Metric

Standard K-Means uses **Euclidean distance**:

$$\text{d}(x, \mu) = \sqrt{\sum_{j=1}^d (x_j - \mu_j)^2}$$

---

## 2.3 Why Mean Minimises Squared Distance

For a cluster  $C$ , centroid is:

$$\mu = \frac{1}{|C|} \sum_{i \in C} x_i$$

This minimises:

$$\sum_{i \in C} |x_i - \mu|^2$$

(by taking derivative and setting to zero)

---

## 3. K-Means Algorithm

### Step-by-Step Procedure

1. Choose K
  2. Initialise K centroids (random or k-means++)
  3. Assign each point to nearest centroid
  4. Update centroids
  5. Repeat until convergence
- 

### Convergence Criteria

- No change in assignments
  - Centroids stop moving
  - Maximum iterations reached
- 

## 4. Initialization Methods

---

## **Random Initialization**

- Fast
  - Can lead to poor local minima
- 

## **K-Means++ (Recommended)**

Selects spread-out initial centroids:

$$\text{P}(x) = \frac{\sum D(x)^2}{\sum D(x)^2}$$

Improves convergence and stability

---

## **5. Choosing the Number of Clusters (K)**

### **Elbow Method**

Plot WCSS vs K and look for elbow

---

### **Silhouette Score**

$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}$$

- $a(i)$ : mean intra-cluster distance
  - $b(i)$ : mean nearest-cluster distance
- 

## **6. Model Evaluation**

### **Internal Metrics**

- Inertia (WCSS)
  - Silhouette Score
  - Davies–Bouldin Index
- 

### **External Metrics (If labels available)**

- Adjusted Rand Index (ARI)
- Normalised Mutual Information (NMI)

---

## 7. Interpretation of Clusters

---

- Centroid coordinates describe typical member
  - Distance to centroid measures cluster fit
  - Clusters are spherical in nature
- 

## 8. Assumptions

- Clusters are spherical
  - Equal cluster variance
  - Similar cluster sizes
  - Features are scaled
- 

## 9. Common Pitfalls & Misconceptions

- ✗ Not scaling features
  - ✗ Assuming K-Means finds global optimum
  - ✗ Using K-Means for non-spherical clusters
  - ✗ Interpreting clusters as ground truth
- 

## 10. Python Implementation

---

### 10.1 From Scratch (NumPy)

```
import numpy as np

class KMeans:
    def __init__(self, k=3, max_iters=100):
        self.k = k
        self.max_iters = max_iters

    def fit(self, X):
        n, d = X.shape
        self.centroids = X[np.random.choice(n, self.k, replace=False)]

        for _ in range(self.max_iters):
            clusters = [[] for _ in range(self.k)]
```

```

        for x in X:
            distances = [np.linalg.norm(x - c) for c in self.centroids]
            clusters[np.argmin(distances)].append(x)

        new_centroids = np.array([np.mean(c, axis=0) for c in clusters])
        if np.allclose(self.centroids, new_centroids):
            break
        self.centroids = new_centroids
    
```

## 10.2 Using scikit-learn

```

from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import silhouette_score

X_scaled = StandardScaler().fit_transform(X)

model = KMeans(n_clusters=3, init='k-means++', n_init=10)
labels = model.fit_predict(X_scaled)

print("Silhouette Score:", silhouette_score(X_scaled, labels))
    
```

## 10.3 Visualization

```

import matplotlib.pyplot as plt

plt.scatter(X_scaled[:,0], X_scaled[:,1], c=labels)
plt.scatter(model.cluster_centers_[:,0], model.cluster_centers_[:,1],
            c='red', marker='x')
plt.title("K-Means Clustering")
plt.show()
    
```

## 11. Advanced Topics

- Mini-batch K-Means
- K-Medoids
- Kernel K-Means

- Spectral clustering comparison
- 

## 12. When NOT to Use K-Means

- Non-spherical clusters
  - Heavy noise/outliers
  - Different cluster densities
- 

## 13. Best Practices

- Always scale features
  - Use k-means++
  - Run multiple initializations
  - Validate with silhouette score
- 

## 14. Summary

K-Means is: - Simple - Fast - Scalable - Geometry-based

But: - Sensitive to initialization - Assumes spherical clusters

---

End of Revision Guide